

Minimum Spanning Trees

CS 4231, Fall 2020

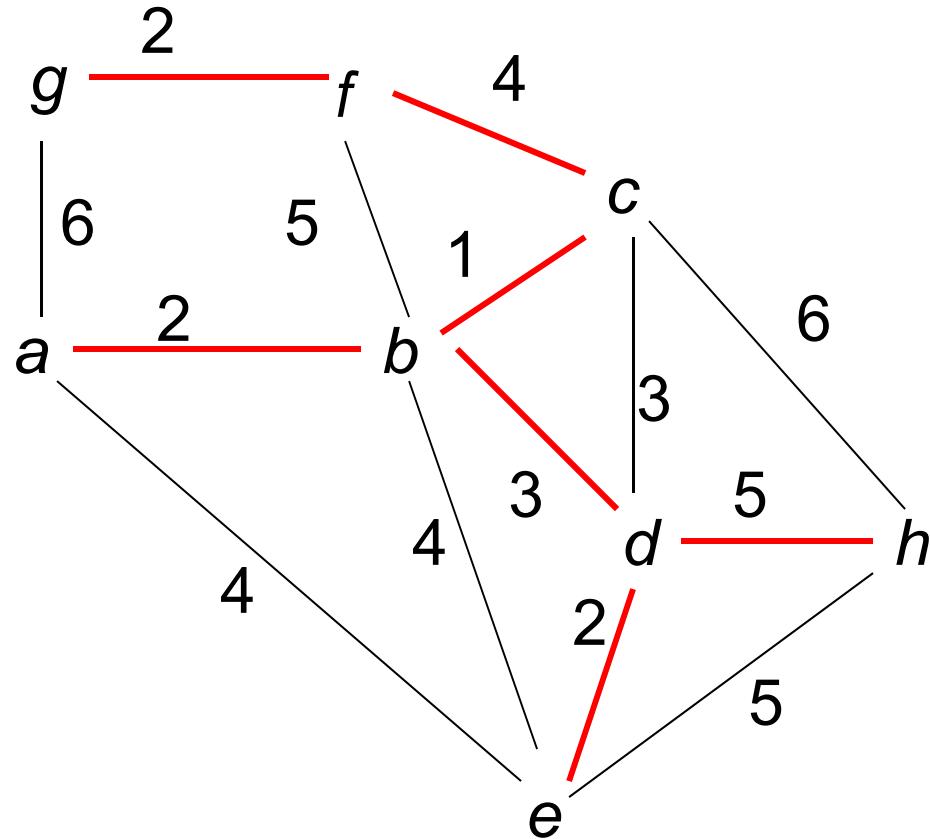
Mihalis Yannakakis

Minimum Spanning Tree

- **Input:** *Undirected*, connected weighted graph
 $G(N,E)$, weights $w: E \rightarrow \mathbb{R}$ (≥ 0 or < 0)
- **Output:** Spanning tree T of minimum weight
(= Min weight subgraph that connects all nodes if weights ≥ 0
Proof: If Cycle \rightarrow remove any edge of cycle)
- **Applications:** Power distribution network (Boruvka 1926),
road-, phone-, TV cable-, computer- network

Maximum Spanning Tree \Leftrightarrow Minimum Spanning Tree
- negate weights

Example



Weight of MST = 19

Minimum Spanning Tree - Algorithms

- **Exhaustive:** Enumerate all spanning trees
- Too many, in general exponential number
- Complete graph: #spanning trees = n^{n-2} (Caley)

Greedy algorithms: Prim, Kruskal

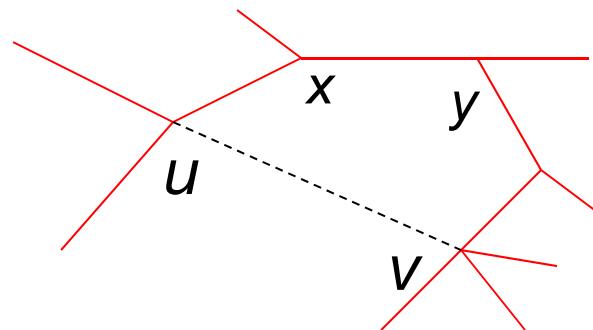
At each step choose minimum weight edge that satisfies a certain criterion

A general underlying mathematical structure (“**matroids**”) in Chapter 16

- analysis of problem => structure of optimal solutions

Exchange argument

Tree T :



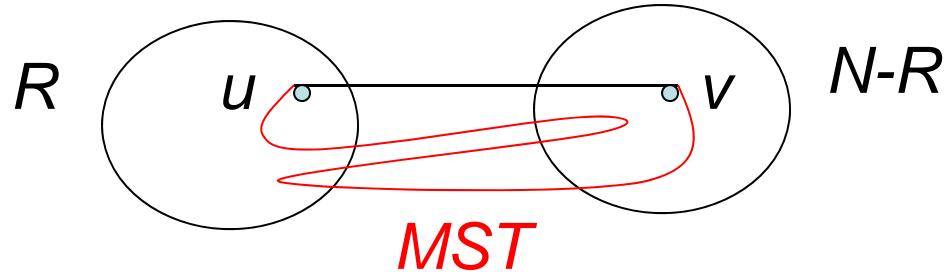
- Given a spanning tree T , adding an out-of-tree edge (u,v) forms a unique cycle: edge (u,v) + $u-v$ path of tree.
 - Exchanging any edge (x,y) of the $u-v$ path with edge (u,v) → another spanning tree T' . Cost lower if $w(u,v) < w(x,y)$
- T is a MST $\Rightarrow \forall$ out-of-tree edge (u,v) has weight \geq maximum weight of the edges along the $u-v$ path in tree T
- Converse also true: HW exercise

Edges across partitions

Partition Theorem: For any set of edges A contained in some MST and for every partition $(R, N-R)$ of the nodes such that no edge of A crosses the partition:

1. Every MST contains some min weight edge across partition (i.e. edge from R to N-R)
2. Every min weight edge across partition \in some MST that contains also A

Proof: Exchange argument for both parts



If $w(u,v) \leq$ weight of MST edges across partition, replace one of these MST edges on $u-v$ path with (u,v)

Some Optimality Conditions

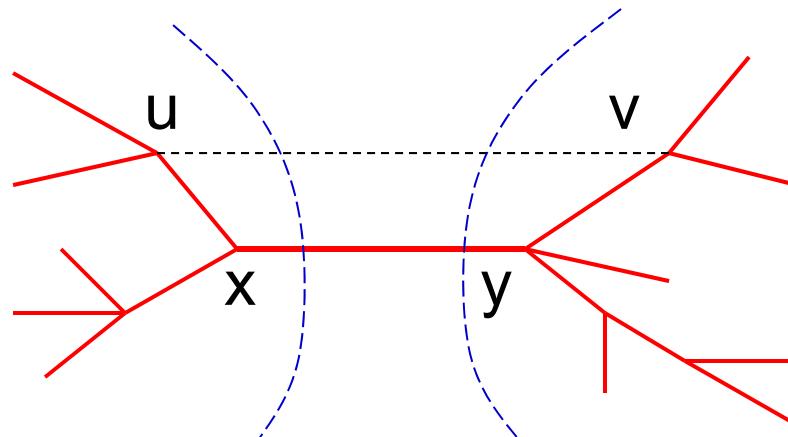
1. A spanning tree T has minimum weight

$$\Leftrightarrow$$

2. \forall out-of-tree edge (u,v) has weight \geq maximum weight of the edges along the $u-v$ path in tree T

$$\Leftrightarrow$$

3. \forall tree edge (x,y) has minimum weight among edges connecting the two components formed by removing the edge from T



Properties

- MST depends only on relative order of edge weights, not their actual values
- If edge weights are distinct \Rightarrow unique MST
- Trick: Can perturb edge weights slightly to break any ties to prove optimality of a spanning tree (simplifies proofs)

HW Exercise:

Prove the optimality conditions and the properties

Prim's algorithm

- Graph Search algorithm from a source node s
Start with $R=\{s\}$ and iteratively add nodes to R .
Policy: in each iteration choose an edge from R to $N-R$ that has minimum weight
- Can implement using a priority queue Q for nodes in $N-R$ with priority of a node = min weight of an edge from a node in R
- Correctness follows from Partition Theorem
Proof by induction on #selected edges that there exists an MST that contains all the edges selected so far.

Prim's Algorithm

MST-Prim(G, w, s)

for each $v \in N$ do $\{d[v] = \infty; p[v] = \perp; \text{mark}[v] = 0\}$

$d[s] = 0;$

$Q = N$ [Q a priority queue with priority $d[]$]

[alternatively, $Q = \{s\}$ initially and insert nodes when reached]

while $Q \neq \emptyset$ do

{ $u = \text{Extract-Min}(Q)$ [Extract-Min operation; first time, $u = s$]

$\text{mark}[u] = 1$

for each $v \in \text{Adj}[u]$ do

if $\text{mark}[v] = 0$ and $d[v] > w(u, v)$ then

$\{d[v] = w(u, v); p[v] = u\}$ [Decrease-Key(v) operation]

}

Worst-Case Time Complexity
Heap: $O(e\log n)$
Fibonacci Heap: $O(e + \log n)$

Time Complexity

Operations:	Extract-Min	Decrease-Key
# of ops:	n	e
Time/Op.		
Heap:	$O(\log n)$	$O(\log n)$
Fibonacci Heap:	$O(\log n)$ <i>(amortized)</i>	$O(1)$

Total (Worst-Case) Time Complexity:

Heap: $O(e \log n)$

Fibonacci Heap: $O(e + n \log n)$

Amortized Analysis of Data Structures

- amortized complexity of operation: bounds the time per operation in any sequence of operations
= average per operation = $(\text{total time}) / \#\text{operations}$
- But no probability: we consider the worst sequence of ops
- An individual operation in the sequence may take lot of time but compensated by earlier cheaper ones
- The Fibonacci heap data structure has amortized complexity $O(\log n)$ for Extract-Min and $O(1)$ for Decrease-Key.

Kruskal's Algorithm

Worst-Case Time Complexity
 $O(e\log n)$

- (1. Sort the edges in nondecreasing order of weight)*
2. Process the edges in order of weight: each edge (u,v) is included in the tree if its nodes u,v are in different connected components in subgraph defined by selected edges so far.

```
for each u in N do Comp(u)={u}
T = ∅
for each edge (u,v) of E in sorted order do
    if ( Comp(u) ≠ Comp(v)) then
        {   T = T ∪ {(u,v)}
            Union sets Comp(u) and Comp(v)
        }
```

*Instead of sorting the edges, we could use a priority queue and extract-min in each iteration, until we have $n-1$ edges in T

Kruskal's Algorithm

1. Sort the edges in nondecreasing order of weight
2. Process the edges in order: each edge (u,v) is included in the tree if its nodes u,v are in different connected components in subgraph defined by selected edges so far.

Correctness: By Partition Theorem

Time Complexity: $O(e \log n)$

Implementation: Need to maintain components

Operations: **Find** component of a node

Union two components

At most $2e$ Find operations (2 for each edge)

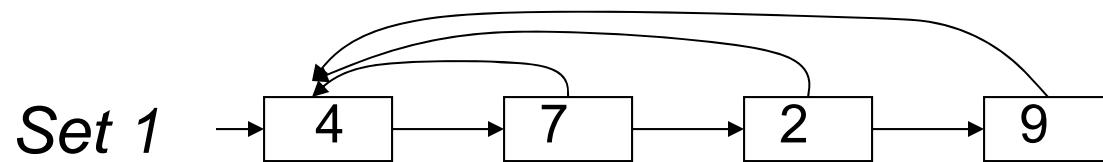
$n-1$ Union operations (1 for each edge of the tree)

Union-Find (Disjoint Sets) Data structure

- Maintain a family of *disjoint sets* over a set N of n elements
- Initially each element in singleton set by itself
- Sequence of operations:
- **FIND(x)**: return (pointer to) set that contains x
- **UNION(S,T)**: union sets S,T

One simple approach: Linked Lists

- Linked list for each set, with pointers from all elements to the head

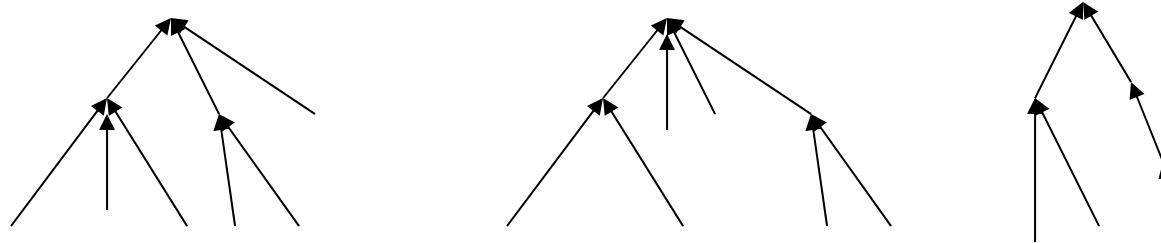


- **FIND:** O(1) time
- **UNION:** Concatenate the lists and update pointers of the elements of one list
- **Key idea:** Update the smaller set
- m FINDS, n UNIONS: Time $O(m + n \log n)$

Reason: Every node has its pointer updated $\leq \log n$ time, because this happens only when the size of its set doubles

Faster: Forest Data structure

- Rooted Tree for each set, with elements at the nodes
(the tree for a set is *not* the same as tree for a component)



FIND(x): Trace path from x to the root of its tree, and hang all nodes of the path from the root (“**path compression**”)

UNION(S,T): Hang the root of one tree (the one with smaller “rank”) from the root of the other

Amortized time per operation: $\alpha(n)$ an extremely slowly growing function: $\alpha(n) \leq 4$ for all realistic n (eg. for $n \leq 10^{80}$)

See Chapter 21 for more details

More advanced MST Algorithms

- If edges sorted \Rightarrow Kruskal almost linear time
- Sorting is *not* required for MST
- Many further algorithms, improving the running time
- Randomized linear time (Karger-Klein-Tarjan'95)
- $O(e \alpha(n))$ deterministic time (Chazelle'2000)
- Optimal time (Petie-Ramachandran'2002)
- Open : Deterministic Linear time?

Applications – Clustering

- Complete Graph
 - Vertices = objects (eg. documents, images, dna seqs...)
 - Edge weights = dissimilarity/distance measure
 - **Clustering problems:** Want similar objects in same cluster, dissimilar in different clusters
 - Various metrics to evaluate clusterings
-
- One metric: Partition objects into k clusters to maximize the minimum distance between any two objects in different clusters
 - Run Kruskal's algorithm till k components
 - HW Exercise: Prove that Kruskal computes the optimal k -clustering under this metric

Applications – Bottleneck paths

- Undirected Graph = network
- edge weights = bandwidth of links
- Bandwidth of path = min bandwidth of an edge on path
- Find max bandwidth path from s to t (or from s to all nodes, or between all pairs)
- Undirected graphs: Maximum Weight Spanning Tree gives max bandwidth paths between all pairs of nodes
 - HW Exercise: prove it